



Performance Anti-Patterns in Hibernate

Patrycja Wegrzynowicz
CTO
Yonita, Inc.

DEVOXX™
the java™ community conference



About Me



- 10+ years of professional experience
- PhD in Computer Science
 - Patterns and anti-patterns, static and dynamic code analysis, language semantics, compiler design
- Active author and speaker
 - JavaOne, Devoxx, JavaZone, Jazoon, OOPSLA, ...
- CTO of Yonita, Inc.
 - Bridge the gap between the industry and the academia
 - Automated defects detection and optimization
 - Performance, security, concurrency, databases
- Twitter @yonlabs

Agenda



- 4 puzzles
 - Short program with curious behavior
 - Question to you (multiple choice)
 - Mystery revealed
 - Lessons learned
- 7 performance anti-patterns
 - Description, consequences, solutions

Disclaimer



I do think Hibernate is a great tool!









Puzzle #1

Less Is More



Who knows CaveatEmptor app?

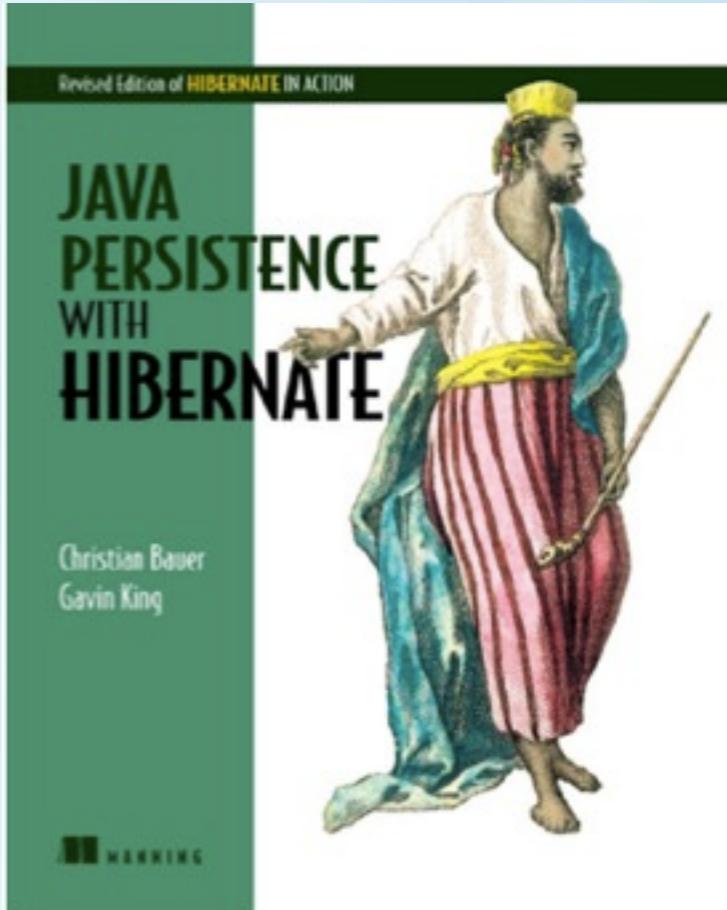


Who's read Java Persistence in Hibernate?

CaveatEmptor



- Demo application from ‘Java Persistence with Hibernate’
- Simple auction system



I'm Gonna Win This Auction!

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
        // ...  
    }  
}
```

What's the Problem?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
        // ...  
    }  
}
```

- (a) None (code is correct)
- (b) Transaction/lock mgmt
- (c) Exception thrown
- (d) None of the above

What's the Problem?



- (a) None (the code is correct)
- (b) Transaction/Lock management**
- (c) Exception thrown
- (d) None of the aboves

The bids can change between calling `getMaxBid` and locking the item.

Another Look

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
        // here, the bids can be modified by a different transaction  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
        // ...  
    }  
}
```

2 Users and Item with 2 Bids



Thread A (new amount 100)

(1)

```
// begin transaction  
// item with 2 bids: 10,20  
Bid curMax = dao.getMaxBid(id);  
// curMax = 20
```

Thread B (new amount 30)

2 Users and Item with 2 Bids



Thread A (new amount 100)

```
(1)  
// begin transaction  
// item with 2 bids: 10,20  
Bid curMax = dao.getMaxBid(id);  
// curMax = 20
```

Thread B (new amount 30)

```
(2)  
// begin transaction  
// item with 2 bids: 10,20  
Bid curMax = dao.getMaxBid(id);  
// curMax = 20
```

2 Users and Item with 2 Bids



Thread A (new amount 100)

```
(1) // begin transaction  
// item with 2 bids: 10,20  
Bid curMax = dao.getMaxBid(id);  
// curMax = 20  
  
(3) // item with 2 bids: 10,20  
// Lock.UPGRADE  
Item item = dao.findById(id,true);  
Bid newBid = item.placeBid(...,  
                           newAmount,  
                           curMax, ...);  
// commit transaction  
// item with 3 bids: 10,20,100
```

Thread B (new amount 30)

```
(2) // begin transaction  
// item with 2 bids: 10,20  
Bid curMax = dao.getMaxBid(id);  
// curMax = 20
```

Successful bid: 100

2 Users and Item with 2 Bids



Thread A (new amount 100)

```
(1) // begin transaction  
// item with 2 bids: 10,20  
Bid curMax = dao.getMaxBid(id);  
// curMax = 20
```

```
(3) // item with 2 bids: 10,20  
// Lock.UPGRADE  
Item item = dao.findById(id,true);  
Bid newBid = item.placeBid(...,  
                           newAmount,  
                           curMax, ...);  
// commit transaction  
// item with 3 bids: 10,20,100
```

Successful bid: 100

Thread B (new amount 30)

```
(2) // begin transaction  
// item with 2 bids: 10,20  
Bid curMax = dao.getMaxBid(id);  
// curMax = 20
```

```
(4) // item with 3 bids: 10,20,100  
// Lock.UPGRADE  
Item item = dao.findById(id,true);  
Bid newBid = item.placeBid(...,  
                           newAmount,  
                           curMax, ...);  
// commit transaction  
// item with 4 bids: 10,20,100,30
```

Successful bid: 30

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
        // ...  
    }  
}
```

How To Fix It?

version and
optimistic locking

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
        // ...  
    }  
}
```

How To Fix It?

version and
optimistic locking

REPEATABLE
READS

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
        // ...  
    }  
}
```

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
        // ...  
    }  
}
```

version and
optimistic locking

REPEATABLE
READS

SERIALIZABLE

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

REPEATABLE
READS

SERIALIZABLE

move lock before
getMinBid/
getMaxBid

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

REPEATABLE
READS

SERIALIZABLE

move lock before
getMinBid/
getMaxBid

different
object model

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

REPEATABLE
READS

SERIALIZABLE

move lock before
getMinBid/
getMaxBid

different
object model

Java
synchronized

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

REPEATABLE
READS

SERIALIZABLE

move lock before
getMinBid/
getMaxBid

different
object model

Java
synchronized

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

NO

REPEATABLE
READS

NO

SERIALIZABLE

move lock before
getMinBid/
getMaxBid

different
object model

Java
synchronized

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

NO

REPEATABLE
READS

NO

SERIYESABLE

move lock before
getMinBid/
getMaxBid

different
object model

Java
synchronized

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

NO

REPEATABLE
READS

NO

SERIYESABLE

move lock before
getBid/
getMaxBid

YES

different
object model

Java
synchronized

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

NO

REPEATABLE
READS

NO

SERIYESABLE

move lock before
getBid/
getMaxBid

YES

different
object model

YES

Java
synchronized

How To Fix It?

```
public class PlaceBidAction implements Action {  
    // invoked in a new transaction  
    public String execute(Map even) {  
        // ...  
        Bid currentMinBid = itemDAO.getMinBid(itemId);  
        Bid currentMaxBid = itemDAO.getMaxBid(itemId);  
  
        // uses load with Lock.UPGRADE  
        Item item = itemDAO.findById(itemId, true);  
        Bid newBid = item.placeBid(userDAO.findById(userId, false),  
            newAmount,  
            currentMaxBid,  
            currentMinBid);  
  
        // ...  
    }  
}
```

version and
optimistic locking

NO

REPEATABLE
READS

NO

SERIYESABLE

move lock before
getBid/
getMaxBid

YES

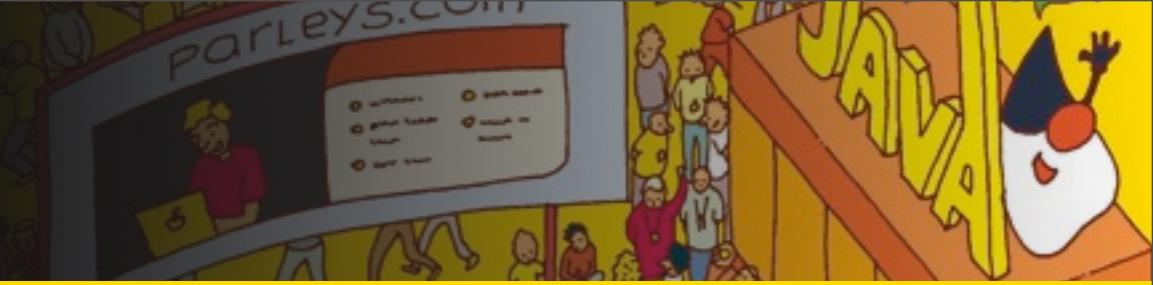
different
object model

YES

leave
synchronized

YES

Lessons Learned



- Nobody's perfect ;-)
- ACID is only a theory
 - Simple transaction interceptor often is not enough
 - Query order does matter
 - Learn more about locks and transaction isolation levels
- And...



Anti-Pattern #1

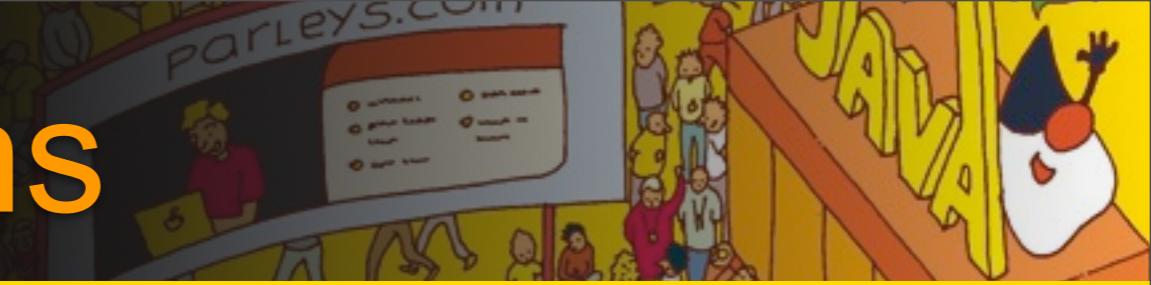
Premature Optimization

Anti-Pattern #1



Premature optimization is the root of all evil

Performance Lessons



- Object model
- Usage patterns
 - Concurrent users
 - Data volume
- Transaction and locking management



Puzzle #2

Heads of Hydra

Heads of Hydra

```
@Entity
public class Hydra {
    private Long id;
    private List<Head> heads = new ArrayList<Head>();

    @Id @GeneratedValue
    public Long getId() {...}
    protected void setId() {...}
    @OneToMany(cascade=CascadeType.ALL)
    public List<Head> getHeads() {
        return Collections.unmodifiableList(heads);
    }
    protected void setHeads(List<Head> heads) {...}
}

// creates and persists the hydra with 3 heads

// new EntityManager and new transaction
Hydra found = em.find(Hydra.class, hydra.getId());
```

How Many Queries in 2nd Tx?

```
@Entity  
public class Hydra {  
    private Long id;  
    private List<Head> heads = new ArrayList<Head>();  
  
    @Id @GeneratedValue  
    public Long getId() {...}  
    protected void setId() {...}  
    @OneToMany(cascade=CascadeType.ALL)  
    public List<Head> getHeads() {  
        return Collections.unmodifiableList(heads);  
    }  
    protected void setHeads(List<Head> heads) {...}  
}  
  
// creates and persists the hydra with 3 heads  
  
// new EntityManager and new transaction  
Hydra found = em.find(Hydra.class, hydra.getId());
```

- (a) 1 select
- (b) 2 selects
- (c) 1+3 selects
- (d) 2 selects, 1 delete, 3 inserts
- (e) None of the above

How Many Queries in 2nd Tx?

- (a) 1 select
- (b) 2 selects
- (c) 1+3 selects
- (d) 2 selects, 1 delete, 3 inserts**
- (e) None of the above

During commit hibernate checks whether the collection property is dirty (needs to be re-created) by comparing Java identities (object references).

Another Look

```
@Entity
public class Hydra {
    private Long id;
    private List<Head> heads = new ArrayList<Head>();

    @Id @GeneratedValue
    public Long getId() {...}
    protected void setId() {...}
    @OneToMany(cascade=CascadeType.ALL)
    public List<Head> getHeads() {
        return Collections.unmodifiableList(heads);
    }
    protected void setHeads(List<Head> heads) {...}
}

// creates and persists hydra with three heads

// new EntityManager and new transaction
// during find only 1 select (hydra)
Hydra found = em.find(Hydra.class, hydra.getId());
// during commit 1 select (heads), 1 delete (heads), 3 inserts (heads)
```

How To Fix It?

```
@Entity
public class Hydra {
    @Id @GeneratedValue
    private Long id;
    @OneToMany(cascade=CascadeType.ALL)
    private List<Head> heads = new ArrayList<Head>();

    public Long getId() {...}
    protected void setId() {...}
    public List<Head> getHeads() {
        return Collections.unmodifiableList(heads);
    }
    protected void setHeads(List<Head> heads) {...}
}

// creates and persists hydra with three heads

// new EntityManager and new transaction
Hydra found = em.find(Hydra.class, hydra.getId());
```

Lessons Learned



- Expect unexpected ;-)
- Prefer field access mappings
- And...



Anti-Pattern #2

Lost Collection Proxy on Owning Side

Lost Collection Proxy

- Description
 - On the owning side of an association, we assign a new collection object to a persistent field/property overriding a collection proxy returned by hibernate
- Consequences
 - The entire collection is re-created: **one delete** to remove all elements and then **many element inserts**

Lost Collection Proxy - Example

```
@Entity
public class Hydra {
    @Id @GeneratedValue
    private Long id;
    @OneToMany(cascade=CascadeType.ALL)
    private List<Head> heads = new ArrayList<Head>();

    public Long getId() {...}
    protected void setId() {...}
    public List<Head> getHeads() {
        return heads;
    }
    public void setHeads(List<Head> heads) {
        this.heads = heads;
    }
}
```

Lost Collection Proxy - Solution

- Operate on collection objects returned by hibernate
 - In most cases it's a much more efficient solution
- Unless you know what you're doing
 - It might be performance-wise in some cases to re-create the entire collection (bulk delete)
 - However, it's the place when hibernate could do better, having all required data available:
 - Original size, added and removed elements
 - $\text{Min}(\#\text{added} + \#\text{removed}, 1 + \#\text{size} - \#\text{removed} + \#\text{added})$



Puzzle #3

Fashionable Developer

Fashionable Developer

```
@Entity
public class Developer {
    @Id @GeneratedValue
    private Long id;
    private String mainTechnology;

    public boolean likesMainTechnology() {
        return "hibernate".equalsIgnoreCase(mainTechnology);
    }
}
// creates and persists a developer that uses hibernate as mainTechnology

// new EntityManager and new transaction
Developer dev = em.find(Developer.class, id);
boolean foundCoolStuff = false;
for (String tech : new String[]{"HTML5", "Android", "Scala"}) {
    dev.setMainTechnology(tech);
    // othersAreUsingIt:
    // select count(*) from Developer where mainTechnology = ? and id != ?
    if (othersAreUsingIt(tech, dev) && dev.likesMainTechnology()) {
        foundCoolStuff = true; break;
    }
}
if (!foundCoolStuff) {
    // still use hibernate
    dev.setMainTechnology("hibernate");
}
```

How Many Queries in 2nd Tx?

```
@Entity
public class Developer {
    @Id @GeneratedValue
    private Long id;
    private String mainTechnology;

    public boolean likesMainTechnology() {
        return "hibernate".equalsIgnoreCase(mainTechnology);
    }
}
// creates and persists a developer that uses hibernate as mainTechnology

// new EntityManager and new transaction
Developer dev = em.find(Developer.class, id);
boolean foundCoolStuff = false;
for (String tech : new String[]{"HTML5", "Android", "Scala"}) {
    dev.setMainTechnology(tech);
    // othersAreUsingIt:
    // select count(*) from Developer where mainTechnology = ? and id != ?
    if (othersAreUsingIt(tech, dev) && dev.likesMainTechnology()) {
        foundCoolStuff = true; break;
    }
}
if (!foundCoolStuff) {
    // still use hibernate
    dev.setMainTechnology("hibernate");
}
```

(a) 2 select
(b) 4 selects
(c) 4 selects, 1 update
(d) 4 selects, 4 updates
(e) None of the above

How Many Queries in 2nd Tx?

- (a) 2 selects
- (b) 4 selects
- (c) 4 selects, 1 update
- (d) 4 selects, 4 inserts**
- (e) None of the above

Hibernate must guarantee correctness of executed queries, therefore in certain situations it must perform flushes during a transaction.

Fashionable Developer

```
@Entity
public class Developer {
    @Id @GeneratedValue
    private Long id;
    private String mainTechnology;

    public boolean likesMainTechnology() {
        return "hibernate".equalsIgnoreCase(mainTechnology);
    }
}
// creates and persists a developer that uses hibernate as mainTechnology

// new EntityManager and new transaction
Developer dev = em.find(Developer.class, id);
boolean foundCoolStuff = false;
for (String tech : new String[]{"HTML5", "Android", "Scala"}) {
    dev.setMainTechnology(tech);
    // othersAreUsingIt:
    // select count(*) from Developer where mainTechnology = ? and id != ?
    if /*othersAreUsingIt(tech, dev) &&*/ dev.likesMainTechnology()) {
        foundCoolStuff = true; break;
    }
}
if (!foundCoolStuff) {
    // still use hibernate
    dev.setMainTechnology("hibernate");
}
```

NO UPDATES



Anti-Pattern #3

Temporary Changes

Temporary Changes

- Description
 - We change the persistent data of an object temporarily to test some features
 - In the meantime a database flush occurs triggered, for example, by a HQL touching the table
- Consequences
 - Flushes can be performed during the transaction, not only at the commit time
 - Temporary changes persisted into the database unnecessarily
- Solution
 - Plan your business model and processing to avoid temporary changes to your persistent data



Puzzle #4

Plant a Tree

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany
    Collection<Tree> trees = new HashSet<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
Tree tree = new Tree("oak");
em.persist(tree);
Forest forest = em.find(Forest.class, id);
forest.plantTree(tree);
```

How Many Queries in 2nd Tx?

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany
    Collection<Tree> trees = new HashSet<Tree>();

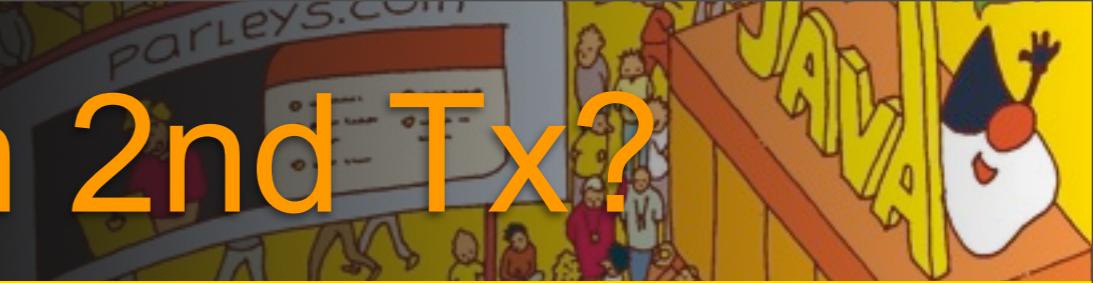
    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
Tree tree = new Tree("oak");
em.persist(tree);
Forest forest = em.find(Forest.class, id);
forest.plantTree(tree);
```

- (a) 1 select, 2 inserts
- (b) 2 selects, 2 inserts
- (c) 2 selects, 1 delete,
10.000+2 inserts
- (d) Even more ;-)

How Many Queries in 2nd Tx?



- (a) 1 select, 2 inserts
- (b) 2 selects, 2 inserts
- (c) 2 selects, 1 delete, 10.000+2 inserts**
- (d) Even more ;-)

The combination of **OneToMany** and **Collection** enables a bag semantic. That's why the collection is re-created.

Plant a Tree - Revisited

```
@Entity
public class Orchard {
    @Id @GeneratedValue
    private Long id;
    @OneToMany
    List<Tree> trees = new ArrayList<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
Tree tree = new Tree("apple tree");
em.persist(tree);
Orchard orchard = em.find(Orchard.class, id);
orchard.plantTree(tree);
```

Plant a Tree - Revisited

```
@Entity
public class Orchard {
    @Id @GeneratedValue
    private Long id;
    @OneToMany
    List<Tree> trees = new ArrayList<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
Tree tree = new Tree("apple tree");
em.persist(tree);
Orchard orchard = em.find(Orchard.class, id);
orchard.plantTree(tree);
```

STILL BAG SEMANTIC!

Plant a Tree - Revisited

```
@Entity  
public class Orchard {  
    @Id @GeneratedValue  
    private Long id;  
    @OneToMany  
    List<Tree> trees = new ArrayList<Tree>();  
  
    public void plantTree(Tree tree) {  
        trees.add(tree);  
    }  
}  
  
// creates and persists a forest with 10.000 trees  
  
// new EntityManager and new transaction  
Tree tree = new Tree("apple tree");  
em.persist(tree);  
Orchard orchard = em.find(Orchard.class, id);  
orchard.plantTree(tree);
```

STILL BAG SEMANTIC!

Use OrderColumn or
IndexColumn for list
semantic.

OneToMany Mapping



Semantic	Java Type	Annotation
bag semantic	java.util.Collection java.util.List	@ElementCollection @OneToMany @ManyToMany
set semantic	java.util.Set	@ElementCollection @OneToMany @ManyToMany
list semantic	java.util.List	(@ElementCollection @OneToMany @ManyToMany) && (@OrderColumn @IndexColumn)

@OneToMany (no cascade option)

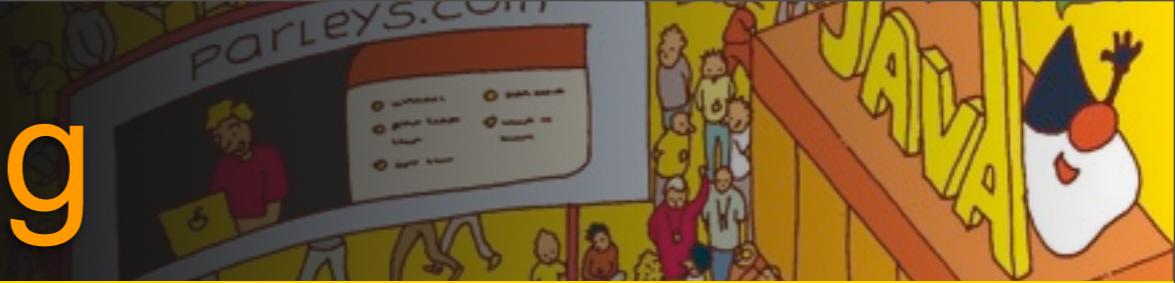
OneToMany Mapping



	add element	remove element	update element
bag semantic	re-create: 1 delete + N inserts	re-create: 1 delete + N inserts	1 update
set semantic	1 insert	1 delete	1 update
list semantic	1 insert + M updates	1 delete + M updates*	1 update

@OneToMany (no cascade option)

OneToMany Mapping



	add element	remove element	update element
bag semantic	re-create: 1 delete + N inserts	re-create: 1 delete + N inserts	1 update
set semantic	1 insert	1 delete	1 update
list semantic	1 insert + M updates	1 delete + M updates*	1 update

@OneToMany (no cascade option)

OneToMany Mapping



	add element	remove element	update element
<p>Oops, we have a problem. (list: removal of nth element, n < size-2)</p>			
Session 1			
Session 2			

```
Hibernate: update Forest_Tree set trees_id=? where Forest_id=? and trees_ORDER=?
2832 [main] WARN org.hibernate.util.JDBCExceptionReporter - SQL Error: 1062, SQLState: 23000
2832 [main] ERROR org.hibernate.util.JDBCExceptionReporter - Duplicate entry '10' for key 'trees_id'
Exception in thread "main" javax.persistence.RollbackException: Error while committing the transaction
    at org.hibernate.ejb.TransactionImpl.commit(TransactionImpl.java:93)
    at com.yonita.examples.jpapuzzles.JPA.execute(JPA.java:23)
    at com.yonita.examples.jpapuzzles.puzzle5.collection.Forest.main(Forest.java:68) <5 internal calls>
Caused by: javax.persistence.PersistenceException: org.hibernate.exception.ConstraintViolationException
    at org.hibernate.ejb.AbstractEntityManagerImpl.convert(AbstractEntityManagerImpl.java:1387)
    at org.hibernate.ejb.AbstractEntityManagerImpl.convert(AbstractEntityManagerImpl.java:1315)
    at org.hibernate.ejb.TransactionImpl.commit(TransactionImpl.java:81)
    ... 7 more
Caused by: org.hibernate.exception.ConstraintViolationException: Could not execute JDBC batch update
```

list
semantic

1 insert + M
updates

1 delete + M
updates*

1 update

@OneToMany (no cascade option)



Anti-Pattern #4

Inadequate Collection Type (Owning Side)

Inadequate Collection Mapping (Owning Side)

- Description
 - Chosen an inadequate collection mapping on the owning side
- Consequences
 - Additional queries executed
- Solution
 - Consider your use cases and usage patterns and choose a mapping with the least performance overhead

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany
    Set<Tree> trees = new HashSet<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
Tree oak = new Tree("oak");
em.persist(oak);
Forest forest = em.find(Forest.class, id);
forest.plantTree(oak);
```



Anti-Pattern #5

OneToMany as Owning Side

OneToMany as Owning Side



- Description
 - OneToMany used as the owning side of a relationship
- Consequences
 - Collection elements loaded into memory
 - Possibly unnecessary queries
 - Transaction and locking schema problems: version, optimistic locking
- Solution
 - Use ManyToOne as the owning side and OneToMany as the inverse side
 - Or at least exclude the collection from locking
 - `@OptimisticLock(excluded=true)`

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany (mappedBy = "forest")
    Set<Tree> trees = new HashSet<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

@Entity
public class Tree {
    @Id @GeneratedValue
    private Long id;
    private String name;
    @ManyToOne
    Forest forest;

    public void setForest(Forest forest) {
        this.forest = forest;
        this.forest.plantTree(this);
    }
}
```

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany (mappedBy = "forest")
    Set<Tree> trees = new HashSet<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

@Entity
public class Tree {
    @Id @GeneratedValue
    private Long id;
    private String name;
    @ManyToOne
    Forest forest;

    public void setForest(Forest forest) {
        this.forest = forest;
        this.forest.plantTree(this);
    }
}
```



Anti-Pattern #6

Inadequate Collection Type (Inverse Side)

Inadequate Collection Type (Inverse Side)

- Description
 - Set used as a collection type on the inverse side of a relationship and to keep the model consistent the owning side updates the inverse side
- Consequences
 - Due to a set semantic, all elements must be loaded when we add a new element
- Solution
 - Use Collection or List instead

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany (mappedBy = "forest")
    Collection<Tree> trees = new HashSet<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

@Entity
public class Tree {
    @Id @GeneratedValue
    private Long id;
    private String name;
    @ManyToOne
    Forest forest;

    public void setForest(Forest forest) {
        this.forest = forest;
        this.forest.plantTree(this);
    }
}
```

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany (mappedBy = "forest")
    Collection<Tree> trees = new HashSet<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}

@Entity
public class Tree {
    @Id @GeneratedValue
    private Long id;
    private String name;
    @ManyToOne
    Forest forest;

    ...
}
// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
em.remove(forest);
```

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany (mappedBy = "forest")
    Collection<Tree> trees = new HashSet<Tree>();

    public void plantTree(Tree tree) {
        trees.add(tree);
    }
}
...
Caused by: java.sql.BatchUpdateException: Cannot delete or update a parent row: a foreign key constraint fails (`jpapuzzles`.`tree`
@E
pu
private Long id;
private String name;
@ManyToOne
Forest forest;

// ...
}
// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
em.remove(forest);
```

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany (mappedBy = "forest")
    Collection<Tree> trees = new HashSet<Tree>();

    // ...
}

@Entity
public class Tree {
    @Id @GeneratedValue
    private Long id;
    private String name;
    @ManyToOne
    Forest forest;

    // ...
}
// creates and persists a forest with 10.000 trees

// new EntityManager and new transaction
for (Tree tree : forest.getTrees()) {
    tree.setForest(null);
}
em.remove(forest);
```

Plant a Tree

```
@Entity
public class Forest {
    @Id @GeneratedValue
    private Long id;
    @OneToMany (mappedBy = "forest")
    Collection<Tree> trees = new HashSet<Tree>();

    // ...
}
```

```
@Entity
public class Tree {
    @Id @GeneratedValue
    private Long id;
    private String name;
    @ManyToOne
    Forest forest;

    // ...
}
// creates and persists a forest with 10.000 trees
```

```
// new EntityManager and new transaction
for (Tree tree : forest.getTrees()) {
    tree.setForest(null);
}
em.remove(forest);
```

10.000 UPDATES
1 DELETE



Anti-Pattern #7

One by One Processing

Inadequate Collection Type (Inverse Side)

- Description
 - Sequential processing of a persistent collection/objects
- Consequences
 - Each database operation performed separately for each element
- Solution
 - Bulk queries
 - That's why CaveatEmptor was implemented like this (aggregate function on a database level)

Conclusions

- Big data cause big problems
 - Standard mappings don't handle large datasets well
 - Smart model, bulk processing, projections
- RTFM
 - Most of information can be found in the hibernate docs
 - Understand what you read
 - Think about consequences
 - And you'll be fine :)
- Hibernate
 - Smarter policies and data structures

Contact

- email: patrycja@yonita.com
- twitter: [@yonlabs](https://twitter.com/@yonlabs)
- www: <http://www.yonita.com>